

Obesity_risk_analysis

October 26, 2025

1 Obesity Risk Analysis and Prediction

```
[1]: import numpy as np
import pandas as pd
from xgboost import XGBClassifier
from sklearn.preprocessing import LabelEncoder, StandardScaler, OneHotEncoder, \
    OrdinalEncoder
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

1.1 Data Loading

```
[2]: train_df = pd.read_csv('train.csv')
test_df = pd.read_csv('test.csv')

print(train_df.head())
print(train_df.columns)
```

	id	Gender	Age	Height	Weight	family_history_with_overweight	\
0	0	Male	24.443011	1.699998	81.669950		yes
1	1	Female	18.000000	1.560000	57.000000		yes
2	2	Female	18.000000	1.711460	50.165754		yes
3	3	Female	20.952737	1.710730	131.274851		yes
4	4	Male	31.641081	1.914186	93.798055		yes

	FAVC	FCVC	NCP	CAEC	SMOKE	CH2O	SCC	FAF	\
0	yes	2.000000	2.983297	Sometimes	no	2.763573	no	0.000000	
1	yes	2.000000	3.000000	Frequently	no	2.000000	no	1.000000	
2	yes	1.880534	1.411685	Sometimes	no	1.910378	no	0.866045	
3	yes	3.000000	3.000000	Sometimes	no	1.674061	no	1.467863	
4	yes	2.679664	1.971472	Sometimes	no	1.979848	no	1.967973	

	TUE	CALC	MTRANS	WeightCategory
0	0.976473	Sometimes	Public_Transportation	Overweight_Level_II
1	1.000000	no	Automobile	Normal_Weight

```

2  1.673584          no  Public_Transportation  Insufficient_Weight
3  0.780199  Sometimes  Public_Transportation    Obesity_Type_III
4  0.931721  Sometimes  Public_Transportation  Overweight_Level_II
Index(['id', 'Gender', 'Age', 'Height', 'Weight',
      'family_history_with_overweight', 'FAVC', 'FCVC', 'NCP', 'CAEC',
      'SMOKE', 'CH2O', 'SCC', 'FAF', 'TUE', 'CALC', 'MTRANS',
      'WeightCategory'],
      dtype='object')

```

```

[3]: print(train_df.info())
      print(train_df.describe())

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15533 entries, 0 to 15532
Data columns (total 18 columns):

```

#	Column	Non-Null Count	Dtype
0	id	15533 non-null	int64
1	Gender	15533 non-null	object
2	Age	15533 non-null	float64
3	Height	15533 non-null	float64
4	Weight	15533 non-null	float64
5	family_history_with_overweight	15533 non-null	object
6	FAVC	15533 non-null	object
7	FCVC	15533 non-null	float64
8	NCP	15533 non-null	float64
9	CAEC	15533 non-null	object
10	SMOKE	15533 non-null	object
11	CH2O	15533 non-null	float64
12	SCC	15533 non-null	object
13	FAF	15533 non-null	float64
14	TUE	15533 non-null	float64
15	CALC	15533 non-null	object
16	MTRANS	15533 non-null	object
17	WeightCategory	15533 non-null	object

dtypes: float64(8), int64(1), object(9)

memory usage: 2.1+ MB

None

	id	Age	Height	Weight	FCVC \
count	15533.000000	15533.000000	15533.000000	15533.000000	15533.000000
mean	7766.000000	23.816308	1.699918	87.785225	2.442917
std	4484.135201	5.663167	0.087670	26.369144	0.530895
min	0.000000	14.000000	1.450000	39.000000	1.000000
25%	3883.000000	20.000000	1.630927	66.000000	2.000000
50%	7766.000000	22.771612	1.700000	84.000000	2.342220
75%	11649.000000	26.000000	1.762921	111.600553	3.000000
max	15532.000000	61.000000	1.975663	165.057269	3.000000

	NCP	CH2O	FAF	TUE
count	15533.000000	15533.000000	15533.000000	15533.000000
mean	2.760425	2.027626	0.976968	0.613813
std	0.706463	0.607733	0.836841	0.602223
min	1.000000	1.000000	0.000000	0.000000
25%	3.000000	1.796257	0.007050	0.000000
50%	3.000000	2.000000	1.000000	0.566353
75%	3.000000	2.531456	1.582675	1.000000
max	4.000000	3.000000	3.000000	2.000000

1.2 Feature Engineering

```
[4]: # -----
# STEP 1: Feature Engineering (Do this first)
# -----
# Create BMI and drop the original columns
for df in [train_df, test_df]:
    # Check for division by zero, although height shouldn't be 0
    df['BMI'] = df['Weight'] / (df['Height'] ** 2)

# Define target and features to drop
target_col = 'WeightCategory'
drop_cols = ['id', 'FCVC', 'FAVC', 'Height', 'Weight'] # Drop originals

numerical_cols = ['Age', 'BMI', 'NCP', 'CH2O', 'FAF', 'TUE']

# Define features for One-Hot Encoding (no order)
nominal_cols = ['MTRANS']

# Define features for Binary Encoding (0/1)
# OrdinalEncoder handles this perfectly.
binary_cols = ['Gender', 'family_history_with_overweight', 'SCC', 'SMOKE']
# Define features for Ordinal Encoding (with order)
caec_order = ['no', 'Sometimes', 'Frequently', 'Always']
calc_order = ['no', 'Sometimes', 'Frequently']
```

1.3 Data Visualization

1.3.1 Finding Outliers for each Feature

```
[5]: # Calculate the IQR for each numeric column
Q1 = train_df[numerical_cols].quantile(0.25)
Q3 = train_df[numerical_cols].quantile(0.75)
IQR = Q3 - Q1

# Define lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
```

```

# Count outliers for each numeric column
outliers_count = ((train_df[numerical_cols] < lower_bound) |
    ↪(train_df[numerical_cols] > upper_bound)).sum()

# Display the count of outliers for each column
print("Count of outliers for each column:")
print(outliers_count)

```

Count of outliers for each column:

```

Age      792
BMI       0
NCP     4548
CH20      0
FAF       0
TUE       0
dtype: int64

```

1.3.2 distribution of various numeric features in the dataset

```

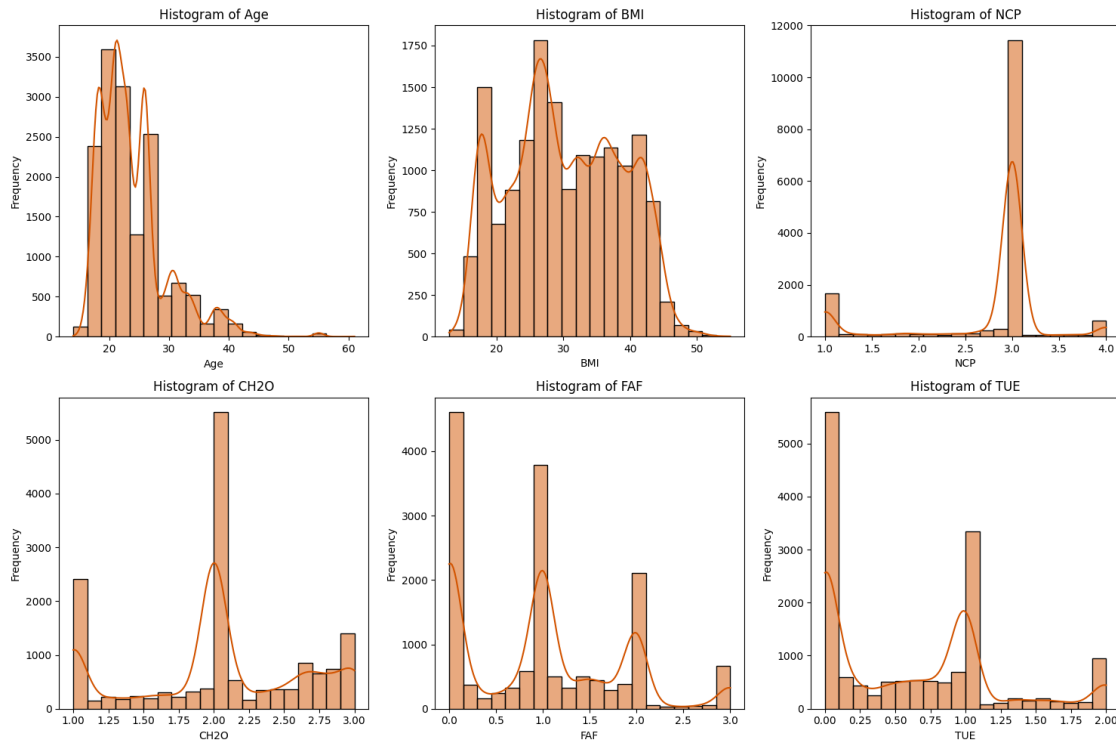
[6]: # Create subplots
num_rows = len(numerical_cols) // 3 + (len(numerical_cols) % 3 > 0)
fig, axes = plt.subplots(num_rows, 3, figsize=(15, 5 * num_rows))
axes = axes.flatten()

# Create histograms for each numeric column
for i, column in enumerate(numerical_cols):
    sns.histplot(train_df[column], kde=True, color='#D35400', bins=20,
    ↪ax=axes[i])
    axes[i].set_title(f'Histogram of {column}')
    axes[i].set_xlabel(column)
    axes[i].set_ylabel('Frequency')

# Hide extra subplots
for j in range(len(numerical_cols), len(axes)):
    axes[j].axis('off')

plt.tight_layout()
plt.show()


```



```
[7]: # Make a fuction that adds data labels for countplots
def add_data_labels(ax):
    """Add data labels to the bars of a countplot."""
    for p in ax.patches:
        ax.annotate(f"{int(p.get_height())}",
                    (p.get_x() + p.get_width() / 2., p.get_height()),
                    ha='center', va='center', fontsize=10, color='black',
                    ↪xytext=(0, 5),
                    textcoords='offset points')
```


1.3.3 distribution of the target variable

```
[8]: # Visualize the distribution of the target variable
plt.figure(figsize=(12, 6))
ax = sns.countplot(x='WeightCategory', data=train_df)
add_data_labels(ax)
plt.title('Distribution of Obesity Levels')
plt.xlabel('Obesity Level')
plt.ylabel('Count')
plt.show()
```




1.3.4 relationship between age and obesity levels

```
[9]: plt.figure(figsize=(12, 6))
sns.boxplot(x='WeightCategory', y='Age', data=train_df)
plt.title('Relationship between Age and Obesity Levels')
plt.xlabel('Obesity Level')
plt.ylabel('Age')
plt.show()
```




1.3.5 Distribution of Age by Gender

```
[10]: plt.figure(figsize=(12, 4))
sns.histplot(data=train_df, x='Age', hue='Gender', kde=True)
# add_data_labels(ax)
plt.title('Distribution of Age by Gender')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```



1.3.6 differences in obesity levels based on family history with overweight

```
[11]: plt.figure(figsize=(8, 6))
ax = sns.countplot(x='family_history_with_overweight', hue='WeightCategory', data=train_df)
add_data_labels(ax)
plt.title('Obesity Levels by Family History with Overweight')
plt.xlabel('Family History with Overweight')
plt.ylabel('Count')
plt.show()
```



1.3.7 Distribution of Weight by Gender

```
[12]: # Create subplots
fig, axes = plt.subplots(1, 2, figsize=(10, 5))



# Plot 1: Boxplot
sns.boxplot(x='Gender', y='Weight', data=train_df, ax=axes[0])
axes[0].set_title('Distribution of Weight by Gender')
axes[0].set_xlabel('Gender')
axes[0].set_ylabel('Weight')
axes[0].set_xticks(ticks=[0, 1])
axes[0].set_xticklabels(['Female', 'Male'])

# Plot 2: Pie chart
train_df['Gender'].value_counts().plot(kind='pie', autopct='%1.1f%%',
    colors=['skyblue', 'lightcoral'], ax=axes[1])
axes[1].set_title('Distribution of Gender')
axes[1].set_ylabel('')
```



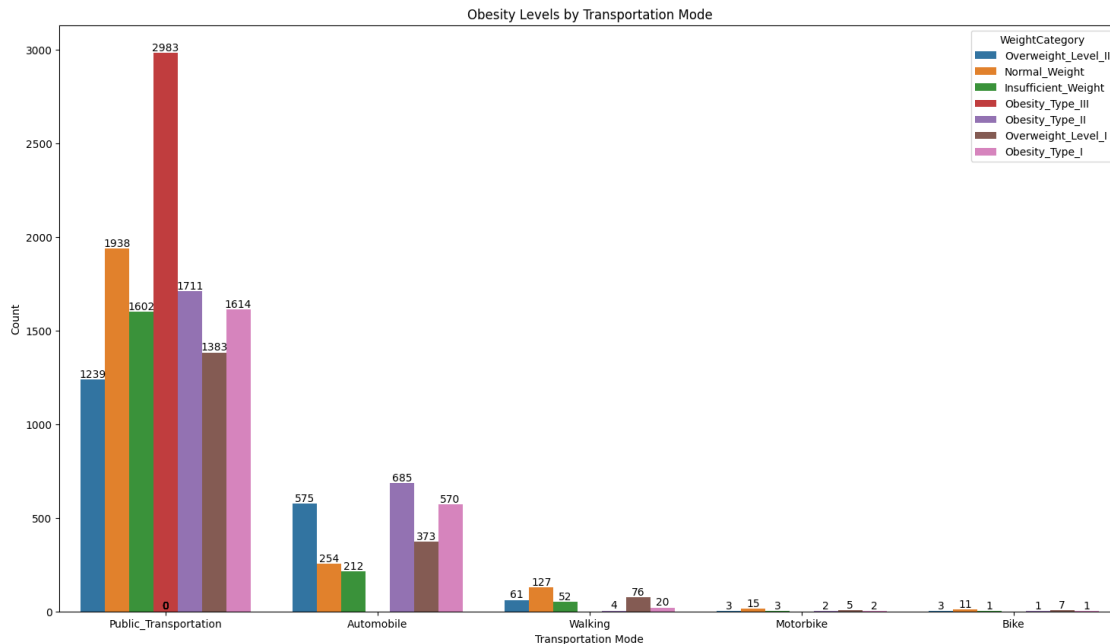
```
# Move legend to bottom center with 2 columns
axes[1].legend(labels=['Female', 'Male'], loc='lower center', bbox_to_anchor=(0.
↪5, -0.1), ncol=2)

plt.tight_layout()
plt.show()
```



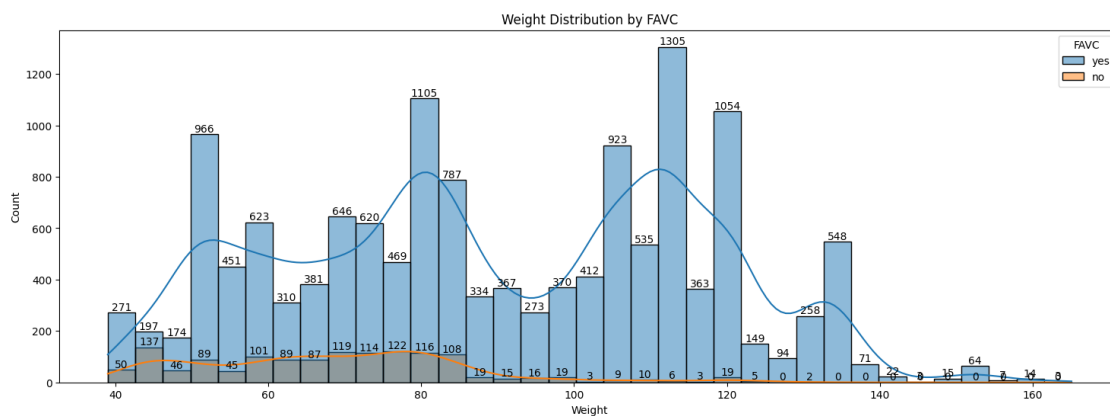
1.3.8 difference in obesity levels based on transportation mode (MTRANS)

```
[13]: plt.figure(figsize=(18, 10))
ax = sns.countplot(x='MTRANS', hue='WeightCategory', data=train_df)
add_data_labels(ax)
plt.title('Obesity Levels by Transportation Mode')
plt.xlabel('Transportation Mode')
plt.ylabel('Count')
plt.show()
```



1.3.9 differences in weight distribution between individuals who frequently consume high-caloric food (FAVC) and those who do not

```
[14]: plt.figure(figsize=(18, 6))
ax = sns.histplot(data=train_df, x='Weight', hue='FAVC', kde=True)
add_data_labels(ax)
plt.title('Weight Distribution by FAVC')
plt.xlabel('Weight')
plt.ylabel('Count')
plt.show()
```



1.4 Data Preprocessing

```
[15]: # -----  
# STEP 3: Create the Preprocessing Pipeline (ColumnTransformer)  
# -----  
  
# This 'preprocessor' is for your features (X)  
preprocessor = ColumnTransformer(  
    transformers=[  
        ('scale', StandardScaler(), numerical_cols),  
        ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False),  
        ↪nominal_cols),  
        ('binary', OrdinalEncoder(), binary_cols),  
        ('ordinal_caec', OrdinalEncoder(categories=[caec_order],  
        ↪handle_unknown='use_encoded_value', unknown_value=-1), ['CAEC']),  
        ('ordinal_calc', OrdinalEncoder(categories=[calc_order],  
        ↪handle_unknown='use_encoded_value', unknown_value=-1), ['CALC'])  
    ],  
    remainder='drop' # Drops any columns we didn't specify (like 'id')  
)  
  
# -----  
# STEP 4: Apply the Pipeline to your Data  
# -----  
  
# Separate features (X) and target (y)  
X_train = train_df.drop(columns=[target_col] + drop_cols, errors='ignore')  
y_train = train_df[target_col]  
  
# The test set from the competition might not have the target column  
if target_col in test_df.columns:  
    X_test = test_df.drop(columns=[target_col] + drop_cols, errors='ignore')  
    y_test = test_df[target_col] # If you have a local validation set  
else:  
    X_test = test_df.drop(columns=drop_cols, errors='ignore')  
  
# Fit the preprocessor on TRAINING data and transform BOTH  
X_train_processed = preprocessor.fit_transform(X_train)  
X_test_processed = preprocessor.transform(X_test)  
  
# -----  
# STEP 5: Encode the Target Variable (y)  
# -----  
  
le2 = LabelEncoder()  
y_train_encoded = le2.fit_transform(y_train)  
feature_names = preprocessor.get_feature_names_out()
```

```

# Convert back to a DataFrame (optional, but good for inspection)
X_train_final = pd.DataFrame(X_train_processed, columns=feature_names)
X_test_final = pd.DataFrame(X_test_processed, columns=feature_names)

print("Preprocessing Complete!")
print("Final training data shape:", X_train_final.shape)
print("Example feature names:", feature_names)

```

Preprocessing Complete!


Final training data shape: (15533, 17)

Example feature names: ['scale__Age' 'scale__BMI' 'scale__NCP' 'scale__CH20'
'scale__FAF'
'scale__TUE' 'onehot__MTRANS_Automobile' 'onehot__MTRANS_Bike'
'onehot__MTRANS_Motorbike' 'onehot__MTRANS_Public_Transportation'
'onehot__MTRANS_Walking' 'binary__Gender'
'binary__family_history_with_overweight' 'binary__SCC' 'binary__SMOKE'
'ordinal_caec__CAEC' 'ordinal_calc__CALC']

```

[16]: # lets visualize correlations
plt.figure(figsize=(12,10))
correlation_matrix = X_train_final.corr()
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap='coolwarm')
plt.title("Correlation Matrix")
plt.show()

```



```
[17]: print(X_train_final.head())
      print(X_test_final.head())
```

	scale__Age	scale__BMI	scale__NCP	scale__CH2O	scale__FAF	scale__TUE \
0	0.110667	-0.235071	0.315486	1.211010	-1.167485	0.602221
1	-1.027075	-0.816880	0.339130	-0.045459	0.027524	0.641290
2	-1.027075	-1.574034	-1.909206	-0.192933	-0.132554	1.759822
3	-0.505665	1.760978	0.339130	-0.581796	0.586624	0.276295
4	1.381740	-0.555043	-1.116801	-0.078619	1.184260	0.527908

	onehot__MTRANS_Automobile	onehot__MTRANS_Bike	onehot__MTRANS_Motorbike \
0	0.0	0.0	0.0
1	1.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

	onehot__MTRANS_Public_Transportation	onehot__MTRANS_Walking	\
0	1.0	0.0	
1	0.0	0.0	
2	1.0	0.0	
3	1.0	0.0	
4	1.0	0.0	

	binary__Gender	binary__family_history_with_overweight	binary__SCC	\
0	1.0	1.0	0.0	
1	0.0	1.0	0.0	
2	0.0	1.0	0.0	
3	0.0	1.0	0.0	
4	1.0	1.0	0.0	

	binary__SMOKE	ordinal_caec__CAEC	ordinal_calc__CALC
0	0.0	1.0	1.0
1	0.0	2.0	0.0
2	0.0	1.0	0.0
3	0.0	1.0	1.0
4	0.0	1.0	1.0

	scale__Age	scale__BMI	scale__NCP	scale__CH20	scale__FAF	scale__TUE	\
0	-0.849222	1.643628	0.339130	-0.033368	0.584315	0.331153	
1	-0.396292	-0.506576	0.339130	-0.045459	0.718028	2.079095	
2	-0.270402	-0.366881	-0.056315	-0.324186	-0.628525	-1.019279	
3	1.253777	0.976798	0.339130	0.192100	-0.315771	-0.853143	
4	-1.027075	-0.830761	0.339130	-0.045459	0.027524	0.641290	

	onehot__MTRANS_Automobile	onehot__MTRANS_Bike	onehot__MTRANS_Motorbike	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	1.0	0.0	0.0	
4	0.0	0.0	0.0	

	onehot__MTRANS_Public_Transportation	onehot__MTRANS_Walking	\
0	1.0	0.0	
1	1.0	0.0	
2	1.0	0.0	
3	0.0	0.0	
4	1.0	0.0	

	binary__Gender	binary__family_history_with_overweight	binary__SCC	\
0	0.0	1.0	0.0	
1	0.0	1.0	0.0	
2	1.0	1.0	0.0	
3	1.0	1.0	0.0	
4	0.0	0.0	0.0	

	binary__SMOKE	ordinal_caec__CAEC	ordinal_calc__CALC
0	0.0	1.0	1.0
1	0.0	1.0	1.0
2	0.0	1.0	1.0
3	0.0	1.0	1.0
4	0.0	1.0	0.0

1.5 Model Definition and Train, Test Splits

```
[18]: XGB = XGBClassifier()

X_train, X_val, y_train, y_val = train_test_split(X_train_final,
↪y_train_encoded, test_size=0.2, random_state=42)
```

```
[19]: # XGB.fit(X_train, y_train)

# # training accuracy
# y_train_pred = XGB.predict(X_train)
# print("Training Predictions:", y_train_pred)
# train_accuracy = accuracy_score(y_train, y_train_pred)
# print(f"Training Accuracy: {train_accuracy*100:.2f}%")

# #validation accuracy
# y_val_pred = XGB.predict(X_val)
# print("Validation Predictions:", y_val_pred)
# val_accuracy = accuracy_score(y_val, y_val_pred)
# print(f"Validation Accuracy: {val_accuracy*100:.2f}%")
```

1.6 HyperParameter Tuning

```
[20]: print("--- Tuning n_estimators and learning_rate ---")

param_grid_1 = {
    'learning_rate': [0.025, 0.05, 0.1, 0.2],
    'n_estimators': [400, 800, 1000, 1200]
}

grid_search_1 = GridSearchCV(estimator=XGB, param_grid=param_grid_1,
                             scoring='accuracy', n_jobs=-1, cv=5, verbose=1)

grid_search_1.fit(X_train, y_train)

print("Best Parameters (Step 1):", grid_search_1.best_params_)
# Store the best parameters to use in the next steps
best_params = grid_search_1.best_params_
```

--- Tuning n_estimators and learning_rate ---

Fitting 5 folds for each of 16 candidates, totalling 80 fits

Best Parameters (Step 1): {'learning_rate': 0.05, 'n_estimators': 400}

```
[21]: print("\n--- Tuning max_depth and min_child_weight ---")

# Update the classifier with the best parameters found so far
xgb_clf_2 = XGBClassifier(
    objective='binary:logistic',
    eval_metric='logloss',
    use_label_encoder=False,
    random_state=42,
    learning_rate=best_params['learning_rate'],
    n_estimators=best_params['n_estimators']
)

param_grid_2 = {
    'max_depth': [3, 4, 5],
    'min_child_weight': [3, 4, 5]
}

grid_search_2 = GridSearchCV(estimator=xgb_clf_2, param_grid=param_grid_2,
                             scoring='accuracy', n_jobs=-1, cv=5, verbose=1)

grid_search_2.fit(X_train, y_train)

print("Best Parameters (Step 2):", grid_search_2.best_params_)
# Update our best parameters dictionary
best_params.update(grid_search_2.best_params_)
```

--- Tuning max_depth and min_child_weight ---

Fitting 5 folds for each of 9 candidates, totalling 45 fits

```
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:
```



```

/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:

```

```

/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:11] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:

```

```

/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:15] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:19] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:19] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:19] WARNING:

```

```

/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:19] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:20] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:20] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:20] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:20] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:20] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:20] WARNING:

```

```

/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:20] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:24] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:24] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:24] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:24] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:24] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:25] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:25] WARNING:

```

```
/workspace/src/learner.cc:738:  
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)  
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-  
packages/xgboost/training.py:183: UserWarning: [22:16:25] WARNING:  
/workspace/src/learner.cc:738:  
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)  
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-  
packages/xgboost/training.py:183: UserWarning: [22:16:25] WARNING:  
/workspace/src/learner.cc:738:  
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)  
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-  
packages/xgboost/training.py:183: UserWarning: [22:16:29] WARNING:  
/workspace/src/learner.cc:738:  
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)  
Best Parameters (Step 2): {'max_depth': 5, 'min_child_weight': 5}
```

```
[22]: print("\n--- Tuning subsample and colsample_bytree ---")  
  
xgb_clf_3 = XGBClassifier(  
    objective='binary:logistic',  
    eval_metric='logloss',  
    use_label_encoder=False,  
    random_state=42,  
    **best_params # Pass all best params found so far  
)  
  
param_grid_3 = {  
    'subsample': [0.8, 0.9, 1.0],  
    'colsample_bytree': [0.9, 1.0]  
}  
  
grid_search_3 = GridSearchCV(estimator=xgb_clf_3, param_grid=param_grid_3,  
                             scoring='accuracy', n_jobs=-1, cv=5, verbose=1)  
  
grid_search_3.fit(X_train, y_train)  
  
print("Best Parameters (Step 3):", grid_search_3.best_params_)  
best_params.update(grid_search_3.best_params_)
```

```

--- Tuning subsample and colsample_bytree ---
Fitting 5 folds for each of 6 candidates, totalling 30 fits

/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

```

Parameters: { "use_label_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:31] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
```


Parameters: { "use_label_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:36] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:37] WARNING:
/workspace/src/learner.cc:738:
```

Parameters: { "use_label_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:42] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:42] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:42] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:42] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:42] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:42] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:45] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
Best Parameters (Step 3): {'colsample_bytree': 1.0, 'subsample': 0.8}
```

```
[23]: print("\n--- Tuning gamma and reg_lambda ---")

xgb_clf_4 = XGBClassifier(
    objective='binary:logistic',
    eval_metric='logloss',
    use_label_encoder=False,
    random_state=42,
    **best_params
)

param_grid_4 = {
    'gamma': [0, 0.1, 0.2, 0.3, 0.4, 0.5],
}

grid_search_4 = GridSearchCV(estimator=xgb_clf_4, param_grid=param_grid_4,
                             scoring='accuracy', n_jobs=-1, cv=5, verbose=1)

grid_search_4.fit(X_train, y_train)

print("Best Parameters (Step 4):", grid_search_4.best_params_)
best_params.update(grid_search_4.best_params_)
```

--- Tuning gamma and reg_lambda ---

Fitting 5 folds for each of 6 candidates, totalling 30 fits

```
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:46] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:52] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:52] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:52] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:52] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:52] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:52] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:52] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:52] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:52] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:53] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:53] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:53] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:57] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:57] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:57] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:58] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:58] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:16:58] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:17:01] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
Best Parameters (Step 4): {'gamma': 0}
```

1.7 Final Model Training

```
[24]: # Create the final, tuned model
final_xgb_clf = XGBClassifier(
    objective='binary:logistic',
    eval_metric='logloss',
    use_label_encoder=False,
    random_state=42,
    **best_params
)

# Train on the full training data
final_xgb_clf.fit(X_train_final, y_train_encoded)
```

```
/home/vansh-doshi/scraping_JD/.venv/lib/python3.12/site-
packages/xgboost/training.py:183: UserWarning: [22:17:03] WARNING:
/workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
[24]: XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=1.0, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric='logloss',
```

```

feature_types=None, feature_weights=None, gamma=0,
grow_policy=None, importance_type=None,
interaction_constraints=None, learning_rate=0.05, max_bin=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=5, max_leaves=None,
min_child_weight=5, missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=400, n_jobs=None,
num_parallel_tree=None, ...)

```

```

[25]: # # training accuracy
# y_train_pred = final_xgb_clf.predict(X_train)
# print("Training Predictions:", y_train_pred)
# train_accuracy = accuracy_score(y_train, y_train_pred)
# print(f"Training Accuracy: {train_accuracy*100:.2f}%")

# #validation accuracy
# y_val_pred = final_xgb_clf.predict(X_val)
# print("Validation Predictions:", y_val_pred)
# val_accuracy = accuracy_score(y_val, y_val_pred)
# print(f"Validation Accuracy: {val_accuracy*100:.2f}%")

```

1.8 Predictions and Making the Submission file for Kaggle Submission

```

[26]: # XGB.fit(X, y)
preds = final_xgb_clf.predict(X_test_final)

```

```

[27]: final_preds = le2.inverse_transform(preds)
print(final_preds)

```

```

['Obesity_Type_III' 'Overweight_Level_I' 'Overweight_Level_II' ...
 'Obesity_Type_I' 'Overweight_Level_II' 'Obesity_Type_II']

```

```

[28]: print(len(final_preds))
print(len(test_df['id']))

```

5225

5225

```

[29]: with open("submission.csv", "w") as f:
    f.write("id,WeightCategory\n")
    for i in range(len(final_preds)):
        f.write(f"{test_df['id'][i]},{final_preds[i]}\n")

```